

# Lab 14: Implement a pre-trained CNN model as a feature extraction using transfer learning

## Aim:

- to use a pre-trained CNN (like VGG16 or ResNet 50)
- as a feature extractor for a new image classifier tool.

## Description:

Transfer learning leverages features learned from large datasets and applies them to smaller, related tasks.

The convolutional base of the model is frozen and new dense layer are added for specific classification outputs.

## Procedure:

- 1) Load a pre-trained model without its top classifier.
- 2) Freeze convolutional layers
- 3) Add a dense layer for the target dataset
- 4) Compile and train the new model.
- 5.) Evaluate accuracy and visualize predictions

## pseudocode:

```
base = VGG16(include_top=False, weights='imagenet')
freeze base layer
Add flatten → Dense → Output layers
compile and train on new dataset
evaluate and predict.
```

## observation:

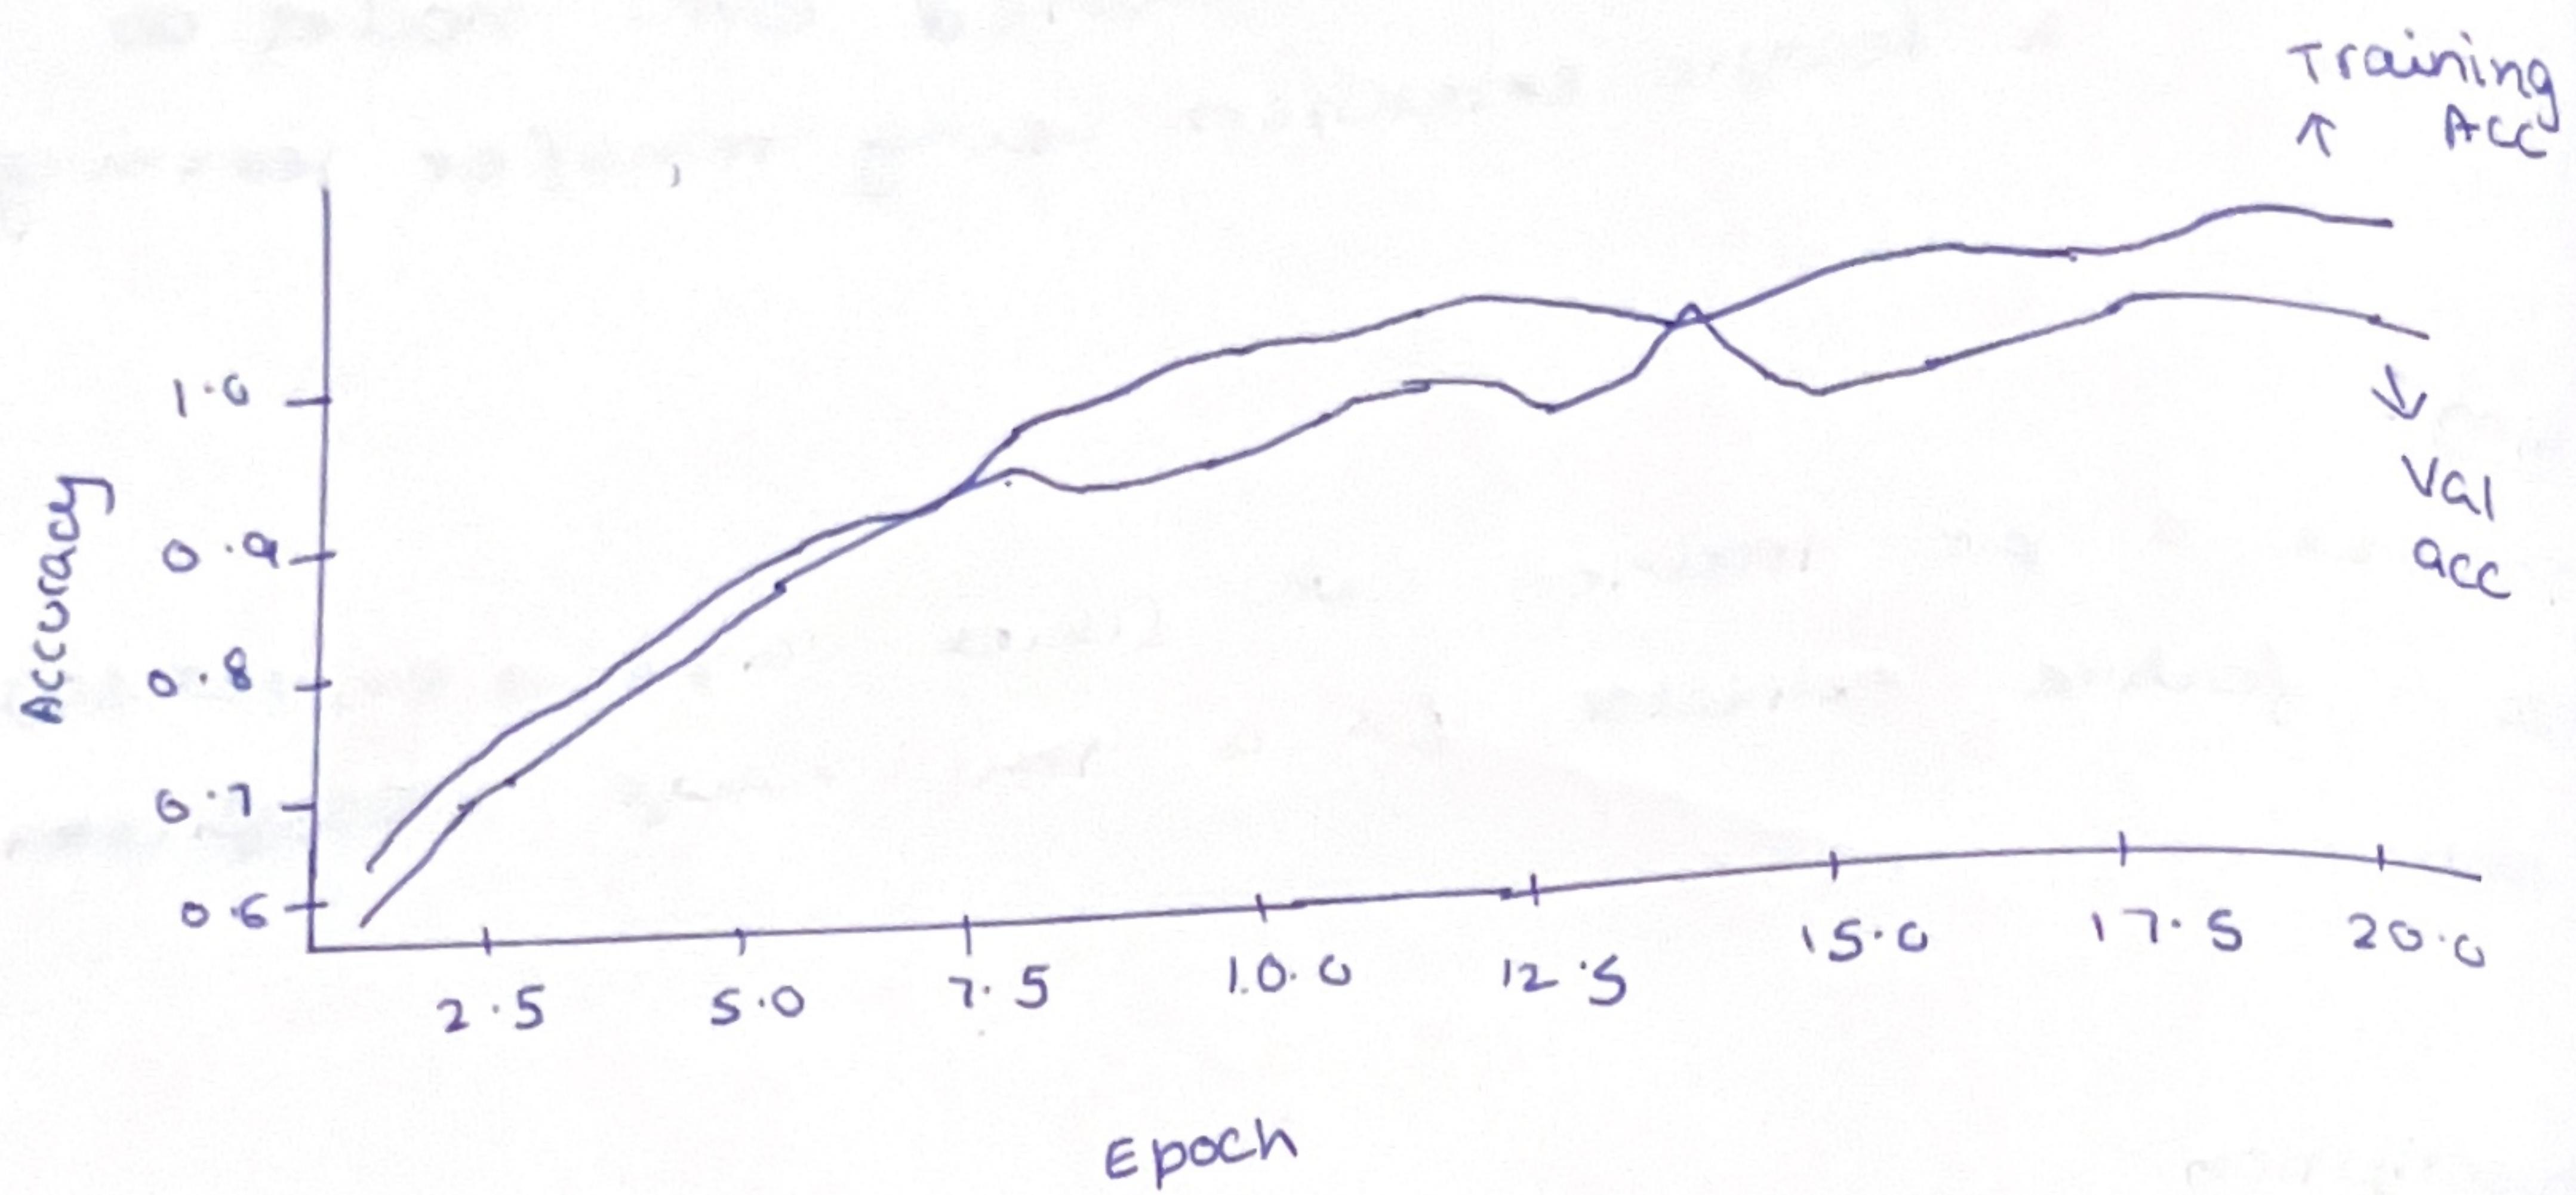
Model achieves high accuracy with fewer epochs due to reused feature maps.

feature extraction reduces training time and data requirements

## result:

successfully used pre trained CNN as a feature extractor for a new classification model.

80%



Epoch	Training Acc	Val. Acc.
1	58.2	55.1
5	72.4	69.3
10	83.7	80.1
15	90.5	87.2
20	99.8	94.7

Week14.ipynb ★ Changes will not be saved

File Edit View Insert Runtime Tools Help

Q Commands + Code + Text ▶ Run all ▾ Copy to Drive

```
( ) ⏪ import torch, torch.nn as nn, torch.optim as optim
      from torchvision import models, datasets, transforms
      from torch.utils.data import DataLoader
      from torchinfo import summary
      from time import time

( ) <> torch.manual_seed(42)
      device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

( ) ⏪ print("Device:", device)

( ) # --- Data (CIFAR-10 -> 224, ImageNet stats)
      img_size = 224
      batch_size = 128
      num_classes = 10

( ) weights18 = models.ResNet18_Weights.IMGNET1K_V1
      mean, std = weights18.meta.get('mean', (0.485, 0.456, 0.406)), weights18.meta.get('std', (0.229, 0.224, 0.225))

( ) tf_train = transforms.Compose([
      transforms.Resize((img_size, img_size)),
      transforms.RandomHorizontalFlip(),
      transforms.RandomCrop(img_size, padding=4),
      transforms.ToTensor(),
      transforms.Normalize(mean=mean, std=std)
])

( ) tf_val = transforms.Compose([
      transforms.Resize((img_size, img_size)),
      transforms.ToTensor(),
      transforms.Normalize(mean=mean, std=std)
```

{} Variables ⌂ Terminal

Q Search

32°C  
Mostly cloudy

```

    x, y = x.to(device), y.to(device)
    logits = model(x)
    loss = criterion(logits, y)

    loss.backward()

    optimizer.step()

    vloss, vacc = evaluate()
    best_acc = max(best_acc, vacc)
    print(f"Epoch {ep}/{epochs} | val_loss={vloss:.4f} | val_acc={vacc*100:.2f}% | (time-{time:.2f}s)")

    print(f"Done. Best val acc (frozen backbone): {best_acc*100:.2f}%")

# --- optional: fine-tune last block + head
for n in [model.layers, model.fc]:
    for p in n.parameters():
        p.requires_grad = True


optimizer = optim.Adam(filter(lambda p: p.requires_grad, model.parameters()), lr=5e-4)

for ep in range(1, epochs_ft+1):
    for x, y in train_loader:
        x, y = x.to(device), y.to(device)
        optimizer.zero_grad()
        loss = criterion(model(x), y)
        loss.backward()
        optimizer.step()

    vloss, vacc = evaluate()
    print(f"Epoch {ep}/{epochs_ft} | val_loss={vloss:.4f} | val_acc={vacc*100:.2f}%")

print("Transfer learning complete.")

```



Q Commands + Code ▾ + Text ▶ Run all ▾ Copy to Drive

Connect ↗

28 Share

```
print("Transfer learning complete.")

Device: cuda
100k|██████████| 170M/170M [00:03<00:00, 48.3MB/s]
downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to /root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
```

Layer (type:depth-idx)

Output Shape	Param #
[1, 10]	--

ResNet	
--------	--

Conv2d: 1-1	[1, 64, 112, 112]
-------------	-------------------

BatchNorm2d: 1-2	[1, 64, 112, 112]
------------------	-------------------

ReLU: 1-3	[1, 64, 112, 112]
-----------	-------------------

MaxPool2d: 1-4	[1, 64, 56, 56]
----------------	-----------------

Sequential: 1-5	[1, 64, 56, 56]
-----------------	-----------------

BasicBlock: 2-1	[1, 64, 56, 56]
-----------------	-----------------

Conv2d: 3-1	[1, 64, 56, 56]
-------------	-----------------

BatchNorm2d: 3-2	[1, 64, 56, 56]
------------------	-----------------

ReLU: 3-3	[1, 64, 56, 56]
-----------	-----------------

Conv2d: 3-4	[1, 64, 56, 56]
-------------	-----------------

BatchNorm2d: 3-5	[1, 64, 56, 56]
------------------	-----------------

ReLU: 3-6	[1, 64, 56, 56]
-----------	-----------------

BasicBlock: 2-2	[1, 64, 56, 56]
-----------------	-----------------

Conv2d: 3-7	[1, 64, 56, 56]
-------------	-----------------

BatchNorm2d: 3-8	(36,864)
------------------	----------

ReLU: 3-9	(128)
-----------	-------

Conv2d: 3-10	--
--------------	----

BatchNorm2d: 3-11	(36,864)
-------------------	----------

ReLU: 3-12	(128)
------------	-------

Sequential: 1-6	--
-----------------	----

BasicBlock: 2-3	--
-----------------	----

Conv2d: 3-13	(73,728)
--------------	----------

Q Commands + Code + Text ▶ Run all ▷ Copy to Drive

```
[BatchNorm2d: 3-47]
[ReLU: 3-48]
[Conv2d: 3-49]
[BatchNorm2d: 3-50]
[ReLU: 3-51]
[AdaptiveAvgPool2d: 1-9]
[Sequential: 1-10]
[Dropout: 2-9]
[Linear: 2-10]
```

```
[1, 512, 7, 7]
[1, 512, 7, 7]
[1, 512, 7, 7]
[1, 512, 7, 7]
[1, 512, 1, 1]
[1, 10]
[1, 512]
[1, 10]
```

5,130

Total params: 11,181,642

Trainable params: 5,130

Non-trainable params: 11,176,512

Total mult-adds (Units.GIGABYTES): 1.81

Input size (MB): 0.60

Forward/backward pass size (MB): 39.74

Params size (MB): 44.73

Estimated Total Size (MB): 85.07

```
Epoch 1/5 | val_loss=0.7023 | val_acc=76.37% | time=105.0s
Epoch 2/5 | val_loss=0.6327 | val_acc=78.72% | time=104.4s
Epoch 3/5 | val_loss=0.6006 | val_acc=79.65% | time=107.7s
Epoch 4/5 | val_loss=0.5957 | val_acc=79.47% | time=103.7s
Epoch 5/5 | val_loss=0.5881 | val_acc=79.87% | time=106.7s
```

Done. Best val Acc (frozen backbone): 79.87%

```
[FT] Epoch 1/3 | val_loss=0.3267 | val_acc=88.66%
[FT] Epoch 2/3 | val_loss=0.2760 | val_acc=91.04%
[FT] Epoch 3/3 | val_loss=0.2866 | val_acc=90.60%
```

Transfer learning complete.