

Aim: To implement gradient descent and backpropagation for training a simple feedforward neural network on the XOR problem

Description:

- * Gradient descent is an optimization algorithm used to minimize the loss function by updating weights in the opposite direction of the gradient.
- * Backpropagation is the process of calculating the gradient of the loss function with respect to each weight using the chain rule, so we can apply gradient descent efficiently.

The steps are:


1. Forward pass - Compute output from inputs
2. Compute loss - difference between predicted and actual values.
3. Backpropagation pass - Compute gradients of loss w.r.t. weights
4. Update weights - use gradient descent rule:

$$w = w - \eta \cdot \frac{dL}{dw}$$

where η is the learning rate.

Procedure:

1. Initialize weights randomly
2. Forward pass
 - compute the hidden layer
 - Compute the final output
3. Compute loss : mean squared error between predicted output and actual output



OUTPUT:


Epoch 0, loss : 0.2558
Epoch 1000, loss : 0.2494
Epoch 2000, loss : 0.2454
Epoch 3000, loss : 0.2047
Epoch 4000, loss : 0.1532
Epoch 5000, loss : 0.1387
Epoch 6000, loss : 0.1336
Epoch 7000, loss : 0.1312
Epoch 8000, loss : 0.1297
Epoch 9000, loss : 0.1288

Final Predictions:

[0.005 300868]
[0.49554213]
[0.95091319]
[0.50319888]

Loss Curve

Loss Curve



4. Backpropagation

5. weight update

update weights and biases using gradient descent

$$w = w - \eta \cdot \frac{dL}{dw}$$

6. Repeat steps 2-5 for several epochs until the loss converges.

7. Print final predictions after training

Result:

The code has been successfully executed and shows the network learned the XOR logic output near 0 for [0, 0] and [1, 1] and near 1 for [0, 1] and [1, 0]

~~effie~~

EXPERIMENT-7


23-9-25

Aim: To build a convolutional neural network model that can classify images of cats and dogs using a labelled dataset.

Description:

A CNN is a deep learning algorithm designed to recognise patterns in visual data. ~~CNN are inspired~~ The core idea is to use a series of convolutions over input to capture features such as edges, textures, and structures.

Architecture -



- Convolution layer - apply filters to images to extract features.
- Activation (ReLU layer) - Adds non-linearity to the output of each convolution.
- Pooling layer - Reduces the spatial dimensions helping reduce computational cost and overfitting.
- F.C layer - After feature extraction, flatten through FC layer to predict the output.
- Softmax func. - Used in output layer to get final layer values in classification.

Pseudocode:

Output

	Accuracy	Loss	val-accuracy
7/10	0.8493	0.9051	0.6909
8/10	0.8497	0.6159	0.7636
9/10	0.8024	0.3842	0.7729
7/10	0.7987	0.4149	0.7445
8/10	0.8501	0.3423	0.7455
6/10	0.8345	0.3607	0.8364
7/10	0.7897	0.4351	0.8041
8/10	0.8389	0.3291	0.7727
9/10	0.8389	0.3139	0.8182
10/10	0.8546	0.3006	0.7909

Workflow of CNN

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
```

```
iris = load_iris()
X = pd.DataFrame(iris.data, columns=iris.feature_names)
y = pd.DataFrame(iris.target, columns=["species"])

print("First 5 rows of dataset:")
print(X.head())


print("\nTarget distribution:")
print(y.value_counts())

# Visualize features
pd.plotting.scatter_matrix(X, figsize=(10, 8), diagonal='kde')
plt.suptitle("Iris Dataset Feature Distributions", fontsize=14)
plt.show()
```


```
First 5 rows of dataset:
   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0            5.1          3.5            1.4           0.2
1            4.9          3.0            1.4           0.2
2            4.7          3.0            1.3           0.2
3            4.6          3.2            1.3           0.2
4            5.0          3.6            1.5           0.2
Target distribution:
species
0      50
1      50
2      50
Name: count, dtype: int64
```

```
File Edit View Insert Runtime Tools Help  
Q Commands + Code + Text ▶ Run all ▾  
[] X = pd.DataFrame(iris.data, columns=iris.feature_names)  
y = pd.DataFrame(iris.target, columns=["species"])  
  
print("First 5 rows of dataset:")  
print(X.head())  
  
print("\nTarget distribution:")  
print(y.value_counts())  
  
# Visualize features  
pd.plotting.scatter_matrix(X, figsize=(10, 8), diagonal='kde')  
plt.suptitle("Iris Dataset Feature Distributions", fontsize=14)  
plt.show()  
  
First 5 rows of dataset:  
   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  
0             5.1          3.5            1.4           0.2  
1             4.9          3.0            1.4           0.2  
2             4.7          3.2            1.3           0.2  
3             4.6          3.1            1.5           0.2  
4             5.0          3.6            1.4           0.2  
  
Target distribution:  
species  
0      50  
1      50  
2      50  
Name: count, dtype: int64
```

Iris Dataset Feature Distributions



Iris Dataset Feature Distributions



File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

```
[ ]    scaler = StandardScaler()
         X_scaled = scaler.fit_transform(X)

        # One-hot encode labels
        encoder = OneHotEncoder(sparse_output=False)
        y_encoded = encoder.fit_transform(y)

        # Train-test split
        X_train, X_test, y_train, y_test = train_test_split(
            X_scaled, y_encoded, test_size=0.2, random_state=42, stratify=y
        )
```

```
[ ] ➜ def sigmoid(z):
         return 1 / (1 + np.exp(-z))

def sigmoid_derivative(a):
    return a * (1 - a)

def softmax(z):
    exp_z = np.exp(z - np.max(z, axis=1, keepdims=True)) # stability
    return exp_z / np.sum(exp_z, axis=1, keepdims=True)

def cross_entropy(y_true, y_pred):
    m = y_true.shape[0]
    eps = 1e-9
    return -np.sum(y_true * np.log(y_pred + eps)) / m

def accuracy(y_true, y_pred):
    return np.mean(np.argmax(y_true, axis=1) == np.argmax(y_pred, axis=1))
```

```
[ ] class SimpleDNN:
    def __init__(self, input_size, hidden_size, output_size, lr=0.01):
        np.random.seed(42)
        self.W1 = np.random.randn(input_size, hidden_size) * 0.01
        self.b1 = np.zeros((1, hidden_size))
        self.W2 = np.random.randn(hidden_size, output_size) * 0.01
        self.b2 = np.zeros((1, output_size))
```

Variables Terminal

```
class SimpleDNN:
    def __init__(self, input_size, hidden_size, output_size, lr=0.01):
        np.random.seed(42)
        self.W1 = np.random.randn(input_size, hidden_size) * 0.01
        self.b1 = np.zeros((1, hidden_size))
        self.W2 = np.random.randn(hidden_size, output_size) * 0.01
        self.b2 = np.zeros((1, output_size))
        self.lr = lr

    def forward(self, X):
        self.Z1 = np.dot(X, self.W1) + self.b1
        self.A1 = sigmoid(self.Z1)
        self.Z2 = np.dot(self.A1, self.W2) + self.b2
        self.A2 = softmax(self.Z2)
        return self.A2

    def backward(self, X, y, output):
        m = X.shape[0]

        # Output layer gradients
        dZ2 = output - y
        dW2 = np.dot(self.A1.T, dZ2) / m
        db2 = np.sum(dZ2, axis=0, keepdims=True) / m

        # Hidden layer gradients
        dA1 = np.dot(dZ2, self.W2.T)
        dZ1 = dA1 * sigmoid_derivative(self.A1)
        dW1 = np.dot(X.T, dZ1) / m
        db1 = np.sum(dZ1, axis=0, keepdims=True) / m

        # Gradient descent update
        self.W1 -= self.lr * dW1
        self.b1 -= self.lr * db1
        self.W2 -= self.lr * dW2
        self.b2 -= self.lr * db2

    def train(self, X, y, epochs=1000):
        history = {"loss": [], "accuracy": []}
        for i in range(epochs):
```

Week-6.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code Run all

```
self.w1 -= self.lr * dw1
self.b1 -= self.lr * db1
self.W2 -= self.lr * dw2
self.b2 -= self.lr * db2

def train(self, X, y, epochs=1000):
    history = {"loss": [], "accuracy": []}
    for i in range(epochs):
        output = self.forward(X)
        loss = cross_entropy(y, output)
        acc = accuracy(y, output)
        self.backward(X, y, output)

        history["loss"].append(loss)
        history["accuracy"].append(acc)

        if i % 100 == 0:
            print(f"Epoch {i}: Loss={loss:.4f}, Accuracy={acc:.4f}")

    return history

def predict(self, X):
    probs = self.forward(X)
    return np.argmax(probs, axis=1)
```

[]

```
dnn = SimpleDNN(input_size=4, hidden_size=8, output_size=3, lr=0.1)
history = dnn.train(X_train, y_train, epochs=1000)
```

Epoch 0: Loss=1.0987, Accuracy=0.3333
Epoch 100: Loss=1.0948, Accuracy=0.8000
Epoch 200: Loss=0.8991, Accuracy=0.7250
Epoch 300: Loss=0.5464, Accuracy=0.8833
Epoch 400: Loss=0.4240, Accuracy=0.9167
Epoch 500: Loss=0.3531, Accuracy=0.9167
Epoch 600: Loss=0.3018, Accuracy=0.9250
Epoch 700: Loss=0.2621, Accuracy=0.9417
Epoch 800: Loss=0.2300, Accuracy=0.9500
Epoch 900: Loss=0.2032, Accuracy=0.9583

[]

```
v pred = dnn.forward(X_test)
```

{ } Variables

Terminal

